



**Samueli**  
School of Engineering

---

# CS-M151B

# Computer Systems Architecture

---

Blaise Tine  
UCLA Computer Science

# Logistics

---

- TA introduction
- Bruincast

# The Power of Abstraction

---

- **What is an abstraction?**
  - Simplifying a complex system using interfaces
  - A fundamental concept in Computer Science
- **Common abstractions we use in computer science?**

# The Power of Abstraction


---

- Reduction of complexity
- Modularity
  - Enhanced focus
  - Reusability
  - Scalability
  - Maintainability
- Portability
- Divergence
  - Separate growth

# The Power of Abstraction

---


- Computer System Abstraction



# The Power of Abstraction

---


- Computer Systems Abstraction



# The Power of Abstraction

---


- The classic OSI Model of Networking



# The Power of Abstraction

---


- Hardware Abstraction



# A basic computer model


---

- Perform tasks
  - Input
    - Instructions
    - Data
  - Execution
  - Output
    - Result



# Early computers

- Operated more like programmable calculators
  - Manual configuration using switches
  - Simple integer arithmetic
  - Instructions physical media
    - Punched cards
    - Electromechanical
  - e.g. Mark I
    - Storage 72 numbers
    - 3 additions / sec




IBM Harvard Mark I (1944)

# The Von-Newman Architecture

---


- First digital computer model (1945)
  - Stored-program concept
    - Electronic program – no punch card
  - Unified Memory structure
    - Instructions and data together
  - CPU Pipeline
    - Fetch->decode->execute



# The Von-Newman Architecture

---


- Basic architecture
  - Input device
  - Central processing unit (CPU)
  - Memory Unit
  - Output Device



# The Von-Newman Architecture

---


- **Memory Unit**
  - Digital storage
  - Store Instructions and data
  - Programmable
  - Flexibility



# The Von-Newman Architecture

---


- Shared address space
  - Address Partitioning
  - Data region
  - Instruction region
  - Inputs/Outputs region



# The Von-Newman Architecture


---

- Shared address space
  - Shared system bus
  - Simple control logic
  - Single memory interface
  - Program counter (PC)
    - Next instruction address




# The Von-Newman Architecture

- **System Bus**
  - Address bus
    - Specify memory address
    - Uni-directional
  - Data bus
    - Specify data to transfer
    - Bi-directional
  - Control bus
    - Schedule transfer
    - Bi-directional – **why?**




# The Von-Newman Architecture

- Limitations
  - Shared bus bottleneck
    - Too many connected devices
    - Control switching
  - One instruction at the time




# The Von-Newman Architecture

- **Sample Exercise:**
  - Assume following system bus design
    - 32-bit address bus
    - 8-bit data bus
    - 2-bit control bus [W][S]
      - W: read/write enable
        - 0 => read
        - 1 => write
      - S: chip select
        - 0 => Memory
        - 1 => I/O




# The Von-Newman Architecture

- Sample Exercise:
  - CPU wants to read a byte from memory at address 0x80000000
  - Step1: CPU places 0x80000000 on address bus.
  - Step2: CPU selects memory for read by placing 00 on control bus
    - W=0, S=0
  - Step3: Memory loads the byte at specified address location in memory and puts it on the data bus
  - Step4: CPU reads the byte from the data bus




# The Von-Newman Architecture

- Sample Exercise:
  - CPU wants to write a byte to I/O at address 0xFF000000
  - Step1: CPU places 0xFF000000 on address bus.
  - Step2: CPU places the byte to be sent on the data bus
  - Step3: CPU selects I/O for write by placing 11 on control bus
    - W=1, S=1
  - Step4: I/O reads byte from the data bus and stores it to the specified address location in I/O



# Harvard Mark-IO Architecture


- Harvard Architecture Model
  - Harvard mark-I (1930-1952)
  - Separate memory buses
    - Instruction memory
    - Data memory



# The Von-Newman Architecture

---

- Harvard Architecture Model
  - Advantages
    - Simultaneous access
      - Fetch data and instructions in parallel
      - Increases performance
    - Can specialize instruction memory
      - Can be read-only (security)
      - Can be non-volatile (flexibility)
      - Can be faster technology (speed)




# The Von-Newman Architecture

---


- Harvard Architecture Model

- Limitations
  - Complexity and cost
    - Multiple memory buses
  - Limited software flexibility
    - Self-modifying code



# The Von-Newman Architecture

- Von-Newman “bottleneck”
  - High memory access latency
  - High I/O latency
  - High CPU idle time
  - Waiting for memory or I/O



# Attendance

---


- Please Fill the form with today's keyword to register your attendance.



# The Von-Newman Architecture


---

- **Mitigating Memory Access latency**
  - We use deeper caches
  - Worsen by multicore
  - Cache area getting larger than CPU
  - The problem still present today




# The Von-Newman Architecture

- **Mitigations**
  - Caches
    - Instruction cache
    - Data cache
  - On-chip memory
  - Direct Memory Access (DMA)
    - Deferring I/O transfer



# Computer Hardware Model


- System Clock
  - Determine transistors speed
- Control Logic
  - Fetch and schedule instructions
- Arithmetic Logic Unit
  - Execute instructions
- Program Counter (PC)
  - Stores current program location
- Register File
  - Local small and fast storage
- Memory
  - External slow and large storage



# Computer Execution Model

---

- Fetch
  - Fetch next instruction
  - From program counter (PC)
- Decode
  - Extract opcode & operands
- Execute
  - Perform the operation
- Update
  - PC and memory



# The Classic CPU Performance Equation

---

- CPU time = Instruction Count x CPI x Clock Cycle Time
  - CPI = Clock Cycle per Instruction
- $\text{CPU time} = \frac{\text{Instruction Count} \times \text{CPI}}{\text{Clock Rate}}$
- How do you increase performance?
  - Reducing the CPU time

# Reducing CPU Time

---

- $$\text{CPU time} = \frac{\text{Instruction Count} \times \text{CPI}}{\text{Clock Rate}}$$
- Reduce Instruction Count
    - Software, compiler, ISA
  - Reduce CPI
    - Microarchitecture
  - Increase Clock Rate
    - Process technology
    - Microarchitecture efficiency (propagation delay)

# Reducing CPU Time (2)

---

- Example: CPU executing a MatMul workload has baseline performance:
  - Instruction Count: one billion ( $10^9$ ) instructions
  - CPI: 2 clock cycles per instructions
  - Clock Rate: 2 GHz
- New process manufacturing increases the clock rate to 2.5 GHz
- What is the speed-up?
  - Baseline CPU Time =  $\frac{10^9 \times 2}{2 \times 10^9} = 1 \text{ sec}$
  - Improved CPU Time =  $\frac{10^9 \times 2}{2.5 \times 10^9} = 0.8 \text{ sec}$
  - Speed-up =  $\frac{\text{Baseline CPU Time}}{\text{Improved CPU Time}} = \frac{1}{0.25} = 1.25x$


# Writing the program

---

- Simply write machine code
  - Early software
- Disadvantages
  - Hard to write
  - Difficult to debug

Possible advantages?

01010101010101010101




# Writing the program

---

- Hand-written machine code advantages
  - Full h/w control
  - Can write efficient code

01010101010101010101



CPU

# Writing the program

---

- Use a high-level language
  - C++, Fortran, Python
  - Higher abstraction
    - Closer to the problem domain
    - Objects and algorithms

```
void main() {  
    print("Hello world!");  
}
```

C++  
Compiler

```
mov r0 r1  
add r2 r0 r1  
mul r1 r3 r2
```

# Writing the program

---

- **Compiler**
  - Convert to assembly code
  - Enable advanced optimizations
    - Requiring program analysis
  - Difficult to do manually

```
void main() {  
    print("Hello world!");  
}
```

C++  
Compiler

```
mov r0 r1  
add r2 r0 r1  
mul r1 r3 r2
```

# Writing the program

---

- **Assembly code**
  - Textual representation of instructions
  - Depend on target machine architecture

```
void main() {  
    print("Hello world!");  
}
```


C++  
Compiler

```
mov r0 r1  
add r2 r0 r1  
mul r1 r3 r2
```

# Writing the program

---

- **Assembler**
  - Convert assembly code to machine code
  - Machine code contains encoded binary
    - Instructions and data
  - Why is this step separate from the compiler?




# Writing the program

- **Handling Machine code variations**
  - Simplified pseudo instructions
  - Context-specific pseudo instructions
  - Context-specific macro expansion
  - Can depend on the platform or machine

## Example: RISC-V Nop pseudo instruction

- $\text{nop} \Leftrightarrow \text{addi } x0, x0, 0$




# Writing the program

---


- Software portability challenges
  - Want to use a different CPU?
  - Want to use a different OS?

```
void main() {  
    print("Hello world!");  
}
```




# Writing the program

- Software Abstraction layers
  - Application Binary Interface
  - Instruction Set Architecture




# Writing the program

- Application Binary Interface
  - A contract between S/W and OS
  - Binary compatibility
    - Can OS launch application A
    - Can application A interact with application B
  - How?
    - Function calls convention
    - Stack frame organization
    - Binary object format
    - Etc..




# Writing the program

- **Instruction Set Architecture**
  - A contract between S/W and H/W
  - Software/Hardware divergence
    - Can evolve independently
  - Hardware abstraction
    - H/W Simulation



# Writing the program

- **Instruction Set Architecture**
  - S/W view
    - Operations
    - Storage locations
  - Hardware view
    - Instructions
    - Data



# Instruction Set Architecture

---

- **Background**
  - IBM System/360 architecture (1964)
  - First introduction of separation between H/W and S/H
  - Features
    - Complex Instruction Set Computer (CISC)
    - x16 32-bit general-purpose registers
    - x4 64-bit floating-point registers

# Instruction Set Architecture


---

- **Mainstream ISA's**
  - x86/x86-64 (AMD64) ISAs: Intel, AMD
  - ARM ISA: Apple, Samsung, Qualcomm
  - MIPS ISA: Embedded systems, network routers, etc..
  - SPARC ISA: Sun Microsystems
  - RISC-V ISA: SiFive, academia, enthusiasts

# How do you design an ISA?

---

- How fast can we fetch?
- How fast can we decode?
- How fast can we execute?
- How efficient will the hardware be?



# Q&A

---