



Samueli
School of Engineering

CS-M151B

Computer Systems Architecture


Blaise Tine
UCLA Computer Science

Agenda

- Logistics
- ISA

Computer Hardware Model

- Control Logic
 - Fetch and schedule Instructions
- Arithmetic Logic Unit
 - Execute instructions
- Register File
 - Array of registers
 - Fast, small storage
 - Inside the CPU
- Memory
 - Slow, large storage
 - Outside the CPU



The Register File

- Array of registers
- Made of fast flip-flop devices
- Addressing an entry
 - Register index
- Operation
 - CPU read/write into array
 - Providing the index
- Given example
 - 4-entry 32-bit register file
 - r2 contains 0xF0F0F0F0

ID	DATA
0	0x00FF00FF
1	0xAA008800
2	0xF0F0F0F0
3	0x00000000


The Main Memory

- Large array of words
- Made of slower DRAM devices
- Addressing an entry
 - word address
- Operation
 - CPU send read/write requests
 - via the system bus
- Given example
 - 256KB 32-bit memory device
 - Total entries = capacity / word size
 - 0x3 contains 0x00000001

ADDRESS	DATA
0	0x00FF00FF
1	0xAA0088AA
2	0xF0F0F0FF
3	0x00000001
.....
.....
0xFFFF	0x00000000


A Memory Address Space

- A mapping contiguous range of address locations that a computer used to access main memory or other peripherals connected to the computer.
- The computer address bus carries the address to access a connected peripheral.




A Memory Address Space

- Example2:
 - CPU receives load instruction:
 - e.g. LD r0, 0xFFFF
 - Load the value of 0xFFFF into r0
 - It will read the mouse input
 - Then save the value it into register r0



A Memory Address Space

- Example2:
 - CPU receives store instruction:
 - e.g. ST r1, 0xF
 - Store the value of r1 into 0xF
 - It will write r1 into memory
 - At the memory location 0xF



Computer Execution Model

- Basic Instruction Format

- Opcode

- Type of operation
 - e.g. ADD, MUL

- Operands

- Source data to use
 - register identifier (e.g. r0)
 - Or memory address (e.g. 0x8000)



What is included in the ISA?

- Instruction formats
- Opcodes description
- Data types
- Registers
- Addressing modes
- Memory alignment
- Endianness
- Control/status registers
- Others
 - Extensions



ISA



ISA: Instruction Formats

- Enable efficient encoding and decoding
 - Encoding classification
 - Fixed fields
- Examples
 - RISC-V R-Type
 - [op7 | rd | func3 | rs1 | rs2 | funct7]
 - ADD x1, x2, x3
 - RISC-V I-Type
 - [op7 | rd | func3 | rs1 | imm12]
 - ADDI x1, x2, 10
 - **What makes this efficient to decode on hardware?**

Whiteboard: Parallel Decode

ISA: Data Type

- **Data types**
 - Types of data that CPU can directly process
 - Size information as well (32-bit or 64-bit)
- **Common Data Types**
 - Integers
 - Floating-points

ISA: Data Type

- Extended data types
 - Half-precision floats (RISC-V Zfh)
 - Strings (x86)
 - Complex numbers (x86)
 - Vectors (Intel AVX, RISC-V Vector)
 - Matrices (Intel AMX)

ISA: Registers

- Defines operand storage types and sizes
 - 32-bit registers
 - 64-bit registers
 - Floating-point registers
 - Double-precision registers
 - Vectors registers
 - Matrices registers

ISA: Registers

- **Managing Architectural states**
 - **General purpose registers:** hold user data
 - **Program counter:** hold next instruction address (Branches)
 - e.g R15 is explicit in ARM
 - **Status/flag registers:** hold branches or exceptions status
 - rounding mode, division by zero
 - e.g. PSW in x86, CPSR in ARM, CSR in RISC-V

ISA: Addressing Modes

- Addressing modes:
Schemes to encode source operands
 1. Immediate addressing
 2. Register addressing
 3. Direct addressing
 4. Indirect addressing
 5. Displacement addressing
 6. Indexed addressing
 7. Relative addressing
 8. Stack addressing

ISA: Addressing Modes

- **Immediate addressing**
- The operand value is directly embedded within the instruction
- **Example:**
 - ADD r1, r1, #5
- **Pros:**
 - Fast to decode
- **Cons:**
 - Limited range due to instruction size
 - Cannot change operand value at runtime

ISA: Addressing Modes

- **Register addressing**
- The operand is located in a register
- **Example:**
 - ADD r1, r1, r2
- **Pros:**
 - Fast to decode
- **Cons:**
 - Need register file access
 - Register dependencies stall due to data hazards

ISA: Addressing Modes

- **Direct memory addressing**
- The memory operand is directly embedded within the instruction
- **Example:**
 - LD r1, #80000000
- **Pros:**
 - Fast to decode
- **Cons:**
 - Limited range due to instruction size
 - Static address

ISA: Addressing Modes

- **Indirect memory addressing**
- The instruction specifies a register that contains the address of the operand
- **Example:**
 - LD r0, (r1)
- **Pros:**
 - Flexible dynamic address
 - Flexible address range
- **Cons:**
 - Need register file access
 - Register dependencies stall due to data hazards

ISA: Addressing Modes

- **Displacement addressing**
- This combines an address held in a register with an offset value from the instruction to calculate the effective address
- **Example:**
 - LD r0, #5(r1)
- **Pros:**
 - Flexibility for offset-based objects (e.g. structs)
 - Compact S/W translation => less instructions
- **Cons:**
 - Need register file access
 - Register dependencies stall due to data hazards
 - Limited offset range due to instruction size
 - Hardware complexity: extra addition

ISA: Addressing Modes

- **Indexed addressing**
- Similar to displacement addressing but typically used with arrays. A base address is combined with an index from another register.
- **Example:**
 - LD r0, r1(r2)
- **Pros:**
 - Flexibility for array-based indexing (e.g. `widgets[i]`)
 - Compact S/W translation => less instructions
- **Cons:**
 - Need register file access
 - Register dependencies stall due to data hazards (2x)
 - Hardware complexity: extra addition

Whiteboard: Struct/Array addressing

Example:

A point p is located at address 0x400 in memory

You want to access element p->y.

Use direct, indirect, displacement, indexed addressing

```
struct Point {  
    int x; // Offset 0  
    int y; // Offset 4  
};
```

```
struct Point *p;  
int val = p->y;
```

Attendance

- Please Fill the form with today's keyword to register your attendance.



ISA: Addressing Modes

- **Relative addressing**
- Often used in branch instructions, where the operand is a memory address relative to the Program Counter (PC)
- **Example:**
 - $\text{JMP } 4 \Rightarrow$ Jump to memory at $\text{PC} + 4$
- **Pros:**
 - Flexible branching: relocatable program
 - Compact S/W translation \Rightarrow less instructions
- **Cons:**
 - Limited range due to instruction size limit
 - Constraint to PC base address

ISA: Addressing Modes

- **Stack addressing**
- Operands are implicitly located on the stack.
- **Example:**
 - PUSH r0
 - POP r1
- **Pros:**
 - Simplify function calls: call stack
 - Automatic memory management: implicit allocation/deallocation
- **Cons:**
 - Restricted access pattern: top of the stack
 - Limited applications: call stack

ISA: Encoding large Immediate

- **Variable-length instructions**
 - Use multiple fetch operations
 - e.g. loading 0xaaff into r0 using 16-bit instructions (8-bit immediate)
 - Instruction word0 = IMX r0, 0xff
 - Instruction word1 = 0xaa

ISA: Encoding large Immediate (2)

- **Pseudo instructions**
 - Higher-level simplified assembly commands
 - e.g. loading 0xaaff into r0 using 16-bit instructions (8-bit immediate)
 - **LA 0xaaff**
 - MOV r1, 0xaa
 - SHL r1, 8 // shift left by 8 bits
 - OR r1, 0xff
 - LD r0, (r1)

ISA: Encoding large Immediate (3)

- **Fixed offset**
 - Assume a fixed offset to construct the full value
 - For instance, using the upper half of the range
 - e.g. loading 0xaaff into r0 using 16-bit instructions (8-bit immediate)
 - LUI r0, 0xaa // load 0xaa into upper half (0xaa00)
 - OR r0, r0, 0xff

ISA: Memory alignment

- Memory can be described as a contiguous array of bytes
- Restrict memory accesses to addresses that are multiples of some fixed value (usually the size of the data being accessed).
- **e.g.** 4-byte alignment: can only read words at addresses 0, 4, 8, 12, 16, ...

7	6	5	4	3	2	1	0
FF	7	0	EE	EF	CC	1B	A0

ISA: Memory alignment (2)

- **Motivation:**

- When data is aligned, it is easier to access it in a single memory operation
- Affect performance and hardware cost

7	6	5	4	3	2	1	0
FF	7	0	EE	EF	CC	1B	A0

ISA: Memory alignment (3)

- **Example:**
 - Assume a 32-bit CPU wants to support unaligned memory
 - But its fetch unit can only load 4 bytes of word at the time.
 - How will CPU execute LD r0, #1?
 - Answer: send two fetch requests:
 - Read w0: EFCC1BA0
 - Read w1: FF0700EE
 - Combine w0 & w1: EEEFCC1B

7	6	5	4	3	2	1	0
FF	7	0	EE	EF	CC	1B	A0

ISA: Endianness

Endianness refers to the order in which bytes are arranged and accessed in computer memory

- **Big-Endian:** the most significant byte (the "big end") of a word is stored at the smallest memory address and the least significant byte at the highest.
 - e.g. word0 = 0x12345678
- **Pros:** human readable, network-friendly format

7	6	5	4	3	2	1	0
0	0	0	0	78	56	34	12


ISA: Endianness

- **Little-Endian:** the least significant byte (the "little end") of a word is stored at the smallest memory address and the most significant byte at the highest.
 - e.g. word0 = 0x12345678
- **Pros:** arithmetic-friendly format for computation

7	6	5	4	3	2	1	0
0	0	0	0	12	34	56	78


How do you design an ISA?

- Goal?
 - Execute the program as fast as possible
- How?
 - Make all the steps execute fast
 - How do you design the ISA to achieve that?




How do you design an ISA?

- How does the ISA affect instruction Fetch?
 - Instruction size
 - Instruction alignment




How do you design an ISA?

- How does the ISA affect instruction Decoding?
 - Instruction formats
 - Addressing modes



How do you design an ISA?

- How does the ISA affect Execution?
 - Opcode description
 - Registers
 - Data types
 - Addressing modes
 - Memory alignment
 - Endianness



Example: 6502 8-bit CPU

- Usage: popular in the 80's: Apple 1, 2, Nintendo NES, Atari 2600
 - Variable-length instruction: 1-3 bytes
 - 1 byte opcode
 - 1-2 bytes operands
 - 5 Registers: A (accumulator), X (index), Y (index), SP (stack), SR (status)

Accumulator-based architecture

- A reserved register (A) used for intermediate operations.
- e.g. LDA 0x10: load immediate 0x10 into accumulator.
 - ADC 0x3: Add immediate 3 to accumulator
- **Pros/Cons of this design?**

Example: 6502 8-bit CPU

- Accumulator benefits?
 - Programmability
 - Simplify register allocation
 - Manual assembly development
 - Efficiency
 - Fast decode: simple instruction set
 - Fast execute: single register operand

Reference: https://en.wikibooks.org/wiki/6502_Assembly

Q&A
